

JVM Memory Management - Garbage Collection, GC Tools



Azat Satklichov

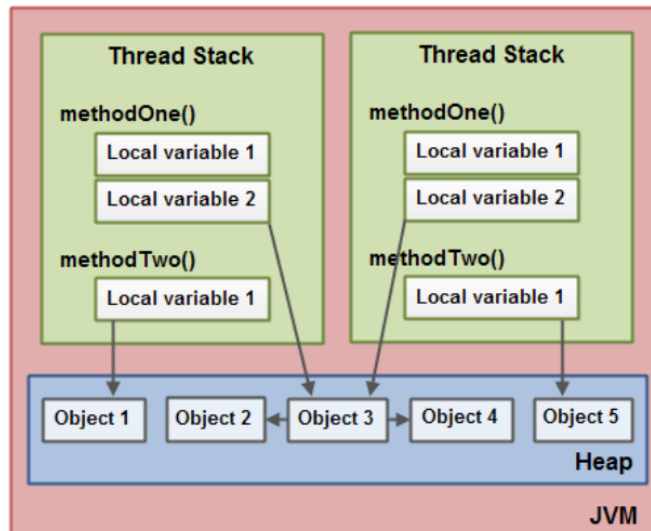
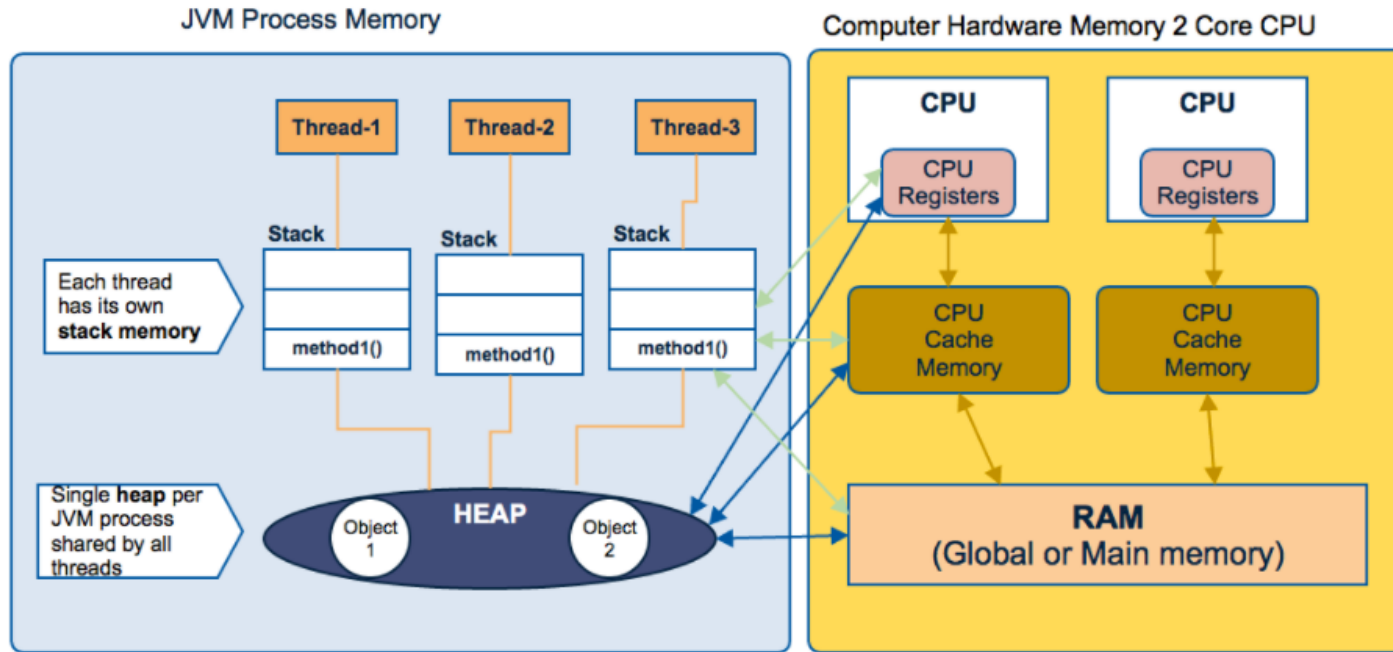
azats@seznam.cz,

<https://github.com/azatsatklichov/Java-Features/tree/master/src/main/java/features/jvm/memory/management>

Agenda

- ❑ Java Garbage Collector
- ❑ Where Storage Lives
- ❑ Types of Garbage Collections
- ❑ Java Memory Model
- ❑ How Garbage Collection Works – Minor, Major and Full GC
- ❑ Garbage Collectors Implementations
- ❑ Garbage Collection Tools - MXBeans, jstat, visualvm, visualgc

Where Storage Lives



- Instance variables and objects live on HEAP
- Methods and Local (automatic/stack/method) vars. live on STACK, but its referencing object or created object live on STACK

Note: local variables created on stack whereas Instance variables created on HEAP

When objects and variables can be stored in various different memory areas in the computer, certain problems may occur. The two main problems are:


- Visibility of thread updates (writes) to shared variables.
- Race conditions when reading, checking and writing shared variables. See [multithreading blog](#)

Hardware memory architecture:

- Registers – fastest storage
- The Stack – in RAM
- Static Storage - (fixed) in RAM
- Constant Storage - ROM
- Non-RAM Storage (persistence)


The **hardware memory architecture does not distinguish between thread stacks and heap**. On the hardware, both the thread stack and the heap are located in main memory.

CODE: WhereStorageLives.java



Java Garbage Collector

Can we force to run GC?



Java's garbage collector provides an automatic solution to memory management.

- Garbage collection cannot ensure that there is enough memory, only it manages available memory as efficiently as possible.
- The garbage collector is under the control of the JVM. The JVM decides when to run the garbage collector.
- JVM will typically run the garbage collector when it senses that memory is running low.
- It tracks each and every object available in the JVM heap space and removes unused ones.
- You can only ask [System.gc()] for GC (can't force) but no guarantee when it happens.
- **Pros:** Automatic memory management (compare by C/C+..) . **Cons:** Managed by JVM. Stop the World issue. Not like man-mem-efficiency
- An object is eligible for garbage collection when can't be reached by any live thread

CODE: Memory Alloc. – GC illustration via addr of Obj using **Unsafe** by Sun.

When Objects are Eligible for Collection?


Once Garbage collector runs, it can discover those objects and delete them.

- Nulling a Reference
- Reassigning a Reference Variable
- Once method call is over BUT ...
- Islands of Isolation //self
- .. ? //pools (interning), caches, ...

```
public static void main(String[] args) {  
    StringBuffer sb = new StringBuffer("Hey GC");  
    System.out.println(sb); //not eligible for collection  
    sb = null; //Nulling a Reference makes it Eligible for GC  
    // Now StringBuffer object is eligible for collection  
}
```

```
public static void main(String[] args) {  
    StringBuffer sb1 = new StringBuffer("Hello GC");  
    StringBuffer sb2 = new StringBuffer("Hi GC");  
    System.out.println(sb1); // "Hello GC" not eligible for collection  
    sb1=sb2; // sb2 re-referenced to "Hi GC" object  
    // Now "Hello GC" object is eligible for collection  
}
```

```
public static void main(String[] args) {  
    IsolatingAReference i2 = new IsolatingAReference();  
    IsolatingAReference i3 = new IsolatingAReference();  
    IsolatingAReference i4 = new IsolatingAReference();  
    i2.i = i3; // i2 refers to i3  
    i3.i = i4; // i3 refers to i4  
    i4.i = i2; // i4 refers to i2  
    i2 = null;  
    i3 = null;  
    i4 = null;  
    // do calc  
}
```




```
public static void main(String[] args) {  
    Date d = getDate();  
    System.out.println("d = " + d);  
}
```

```
public static Date getDate() {  
    Date d2 = new Date();  
    StringBuffer now = new StringBuffer(d2.toString());  
    System.out.println(now);  
    return d2;  
    //StringBuffer object will be eligible for Collection  
    //Date object will not be  
}
```

Types of Garbage Collection


- **Reference Counting** – Count inc. once object started to be referenced, decreased until zero and become eligible for GC. Issue with circular ref. It happens when one object refers to another, and that other one refers to the first object.
- **Mark and Sweep** – “**mark**” phase identifies the objects are still in being used, “**sweep**” phase removes un-used objs. “**compact**” phase defragments the memory.



Types of Garbage Collection

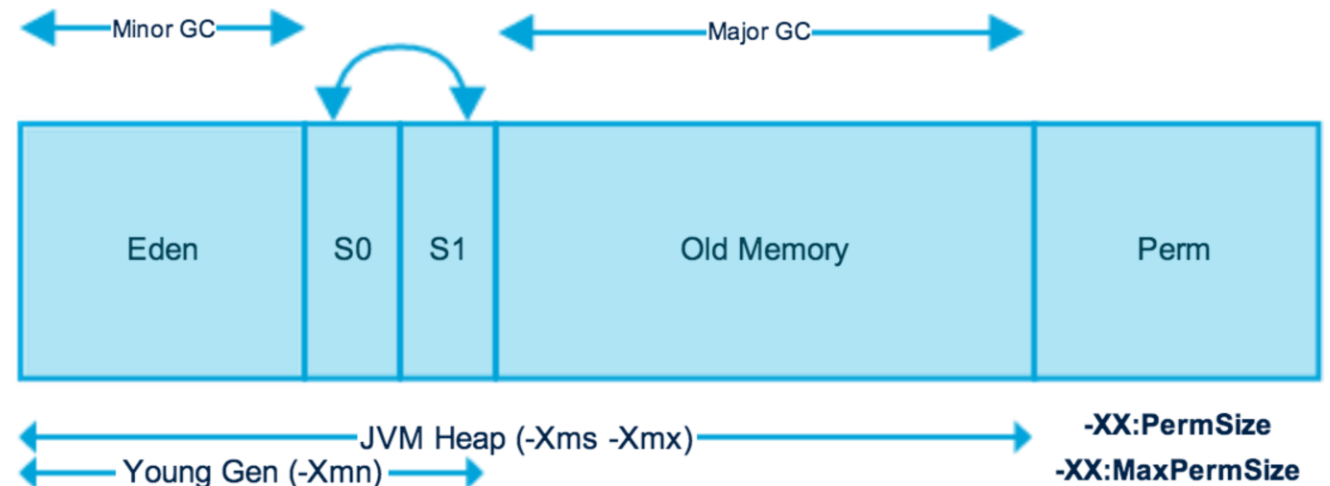
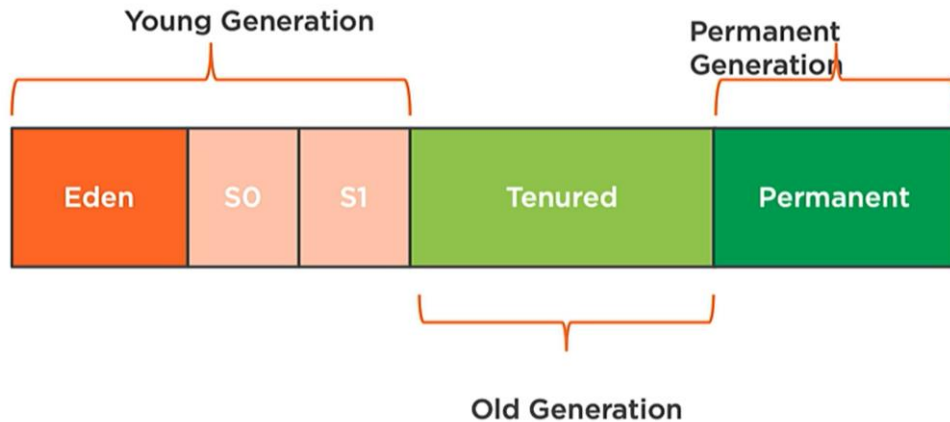
- **Copying** – mark & sweep happens but uses different spaces to manage the memory. During copying memory is compacted at the same time, and from space is cleared after all live objects are moved.
- **Generational** – Manages different generations for memory, long living objects (survived) are promoted to diff-gen. Sweeps happen more often in young generations than old one.
- **Incremental** – does not scan all the memory all the time.

Copying Fromspace to Tospace



Java Memory Model


- Memory space is broken into two spaces: Young and Old (**Tenured Space**) Generations.
 - The JVM internally separates the Young Generation into **Eden** and **Survivor** Spaces (S0, S1).
 - Newly allocated objects are going to Eden [**Eden Gardens**] space.
 - Objects **survived** in young generation promoted to one of Survivor spaces.
 - One survivor space is used at a time, and objects are copied between survivor spaces
 - Old generation where long lived objects (survived many times) live (most stable objects, ..)
 - GC runs frequently on Young Gen. **So motto is: die young or live forever.**
-
- Permanent space (never GC run here) – used by Java Runtime (class info stored,)



How Garbage Collection Works - Minor GC

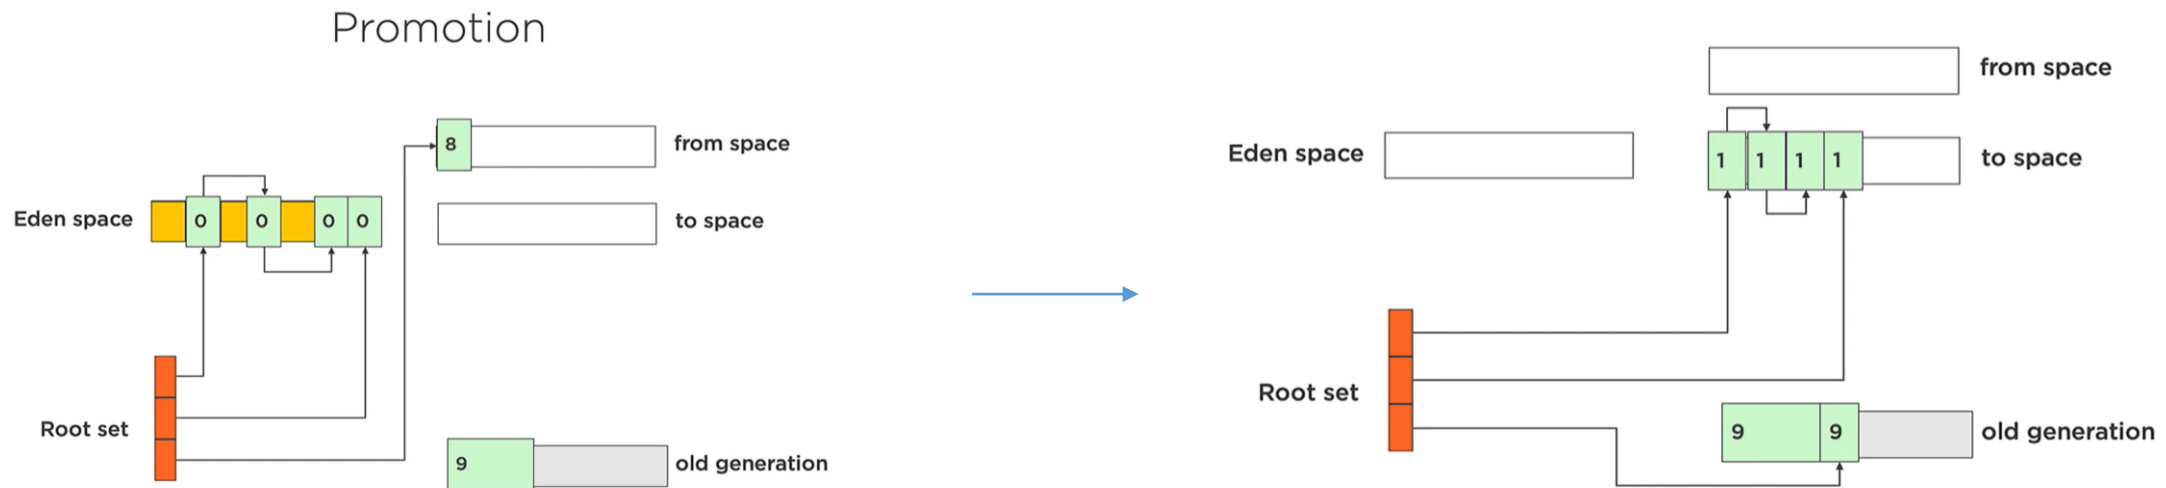
- Minor GC happens when the Young Generation is FULL, happens more frequently.
- New objects are allocated to Eden space. Initial survivor space is called 'from space'
- When GC runs (Minor GC happens), objects are copied to newer survivor space, eden is cleared, and survivor spaces will be swapped

Young Generation Allocation



How Garbage Collection Works – Major GC, Full GC

- **Major GC** triggered when tenured (Old Generation) is FULL. Happens less frequently. It is slower than Minor GC
- **Full GC** - Collects OLD and YOUNG gen. at the same time (on Oracle Java VM)
- In Major CG memory is copied from Young to Old generation – called **promotion**
 - after survived several times
 - once survivor space is full.
 - Or JVM has been told, create object in Old space (via JVM flag set: **-XX:+AlwaysTenure**)



Allocating Objects to Old Space

- Objects with certain size can be directly allocated to Old space. There is no JVM flag for this to allocate always.
- But can be done via: **-XX:PretenureSizeThreshold=<n>** if objSize>n then objects allocated to Old gen. If object size fits to **TLAB**, then JVM allocates it into TLAB (Thread Local Allocation Buffer).

Garbage Collectors Implementations

Which GC to choose? Consider: Stop the world events (1), memory fragmentation(2), Throughput(3). And profile (via `mxbean`, `jstat`, `visualgc`) the app. As close to production load.

- **Serial Garbage Collector** – simplest GC impl., works with a single thread (1) :

>`java -XX:+UseSerialGC -jar App.java`

Good for apps. that do not have small pause time req. and to run on client-style machines.

- **Parallel [Parallel GC, Parallel Old GC] Garbage Collector** – also called Throughput Collector (3). Uses multiple threads for managing heap space. Mostly used in production servers.

Parallel for MinorGC and serial (1) for MajorGC `-XX:+UseParallelGC`, and parallel for both `-XX:+UseParallelOldGC`.

Tune max. gc. *threads and pause time, throughput, and footprint* (heap size). Max-pause-goal: `-XX:MaxGCPauseMillis=<N>`

- **CMS Garbage Collector (low latency)** – Deprecated in Java 9, removed in 14. Concurrent Mark Sweep uses multiple garbage collector threads for gc. designed for apps. in need shorter gc. pauses, throughput is higher ...

`-XX:+UseConcMarkSweepGC`, `-XX:-UseParNewGC`. With serial (1) young space collector

`-XX:+UseConcMarkSweepGC`, `-XX:+UseParNewGC`. With parallel young space collector


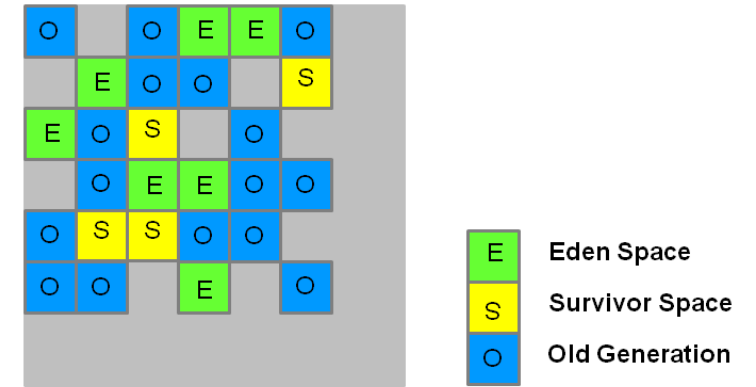
We can limit the number of threads in CMS collector using `-XX:ParallelCMSThreads=<n>` JVM option.

Java 8 JVM has new param – Optimizing heap-memory for reducing the unnecessary use of memory by creating too many instances of the same *String*. > `-XX:+UseStringDeduplication`


Garbage Collectors Implementations

- **G1 (Garbage First, or compacting-collector) Garbage Collector** – incremental GC. Paper 2004, experimental in Java 6, official in Java 7, and default GC since Java 9. Designed for apps running on multi-cores with large memory space. Goal is: **Throughput/Latency balance**. Tuneable pause goals. A bit more CPU intensive. Regions (#2000, size 2MB for 4GB heap)
- Evacuation (moved/copied between regions). Minor & Major collection,...
- `>java -XX:+UseG1GC -jar App.java`

G1 Heap Allocation



- **Epsilon Collector** – A **No-Op GC**. Apps with predictable, bounded memory usage, performance[stress] testing, short-lived obj. Allocates memory but not collect any garbage (memory allocation), once the Java heap is exhausted, the JVM will shut down. Params: -
`XX:+UnlockExperimentalVMOptions, -XX:+UseEpsilonGC`
- **Z Garbage Collector (ZGC is scalable low latency GC)** [pause time **under 10 ms**, maybe in future 1 ms] – is a scalable low-latency garbage collector, experimental in Java 11 for **Linux**, in Java 14 for **Win. & macOS**. Since Java 15 ZGC is onwards. Single generation(all opt. done here via colored pointers), **No pause time increase with heap size increase**, Scale to **multi-terabyte heaps**. -
`XX:+UnlockExperimentalVMOptions, -XX:+UseZGC`
- **Shenandoah Garbage Collector** (derived from G1, scalable low-latency garbage collector) – experimental in Java 12, [Brooks pointers] support large Heaps, Low pause times. Contributed by **RedHat** to OpenJDK. Became product feature in Java 15. Oracle JDK and Open JDK has not this feature. Single generation. Can be used in Java 8 so no need colored pointers like ZGC.
`-XX:+UnlockExperimentalVMOptions -XX:+UseShenandoahGC`



The Z Garbage Collector (ZGC): Low Latency in JDK 11 with Per Liden

At a Glance

- Concurrent

- ✓ Marking
- ✓ Relocation/Compaction
- ✓ Relocation Set Selection
- ✓ Reference Processing
- ✓ JNI WeakRefs Cleaning
- ✓ StringTable Cleaning

- Single generation

- Load barriers
- Colored pointers
- Region-based
- Partial compaction
- NUMA-aware

Garbage Collection Tools

- **MX Beans** – management bean [[GarbageCollectorMXBean](#)] to monitor GC, to get name of garbage collections, numbers of collectors, times of collections and info about memories, ... Each bean for per memory manager (old, young)

CODE: MainMXBeanDemo.java, E_BigObjMemoryAdrrs

Run with: a) `//default GC`

b) `-XX:+UseParallelGC`

c) `-XX:+UseConcMarkSweepGC`

- **Jstat** – command line tool to monitor memory and GC activities, also can be used to profile remote JVM

Run with: `jstat -option <pid> <interval> <count>`

E.g. `>jstat -gc 10632 100 10`

Try with: `-XX:+UseParallelGC`, and `G1`

See jstat descriptions [here](#)

```
C:\Users\as892333>jps
25824 Jps
26576 Jps
30324 E_BigObjMemoryAdrrs
27996 Launcher
33196 Eclipse

C:\Users\as892333>jstat -gcutil 30324
  S0    S1     E      O      M     CCS    YGC     YGCT    FGC     FGCT    CGC     CGCT     GCT
  0.00  99.99  35.99  50.99  79.89  66.89   386    16.563   0.000   0.000   22     0.115    16.678

C:\Users\as892333>jstat -gccause 30324
  S0    S1     E      O      M     CCS    YGC     YGCT    FGC     FGCT    CGC     CGCT     GCT    LGCC     GCC
  0.00  99.99  37.99  52.47  80.33  66.89   512    22.065   0.000   0.000   28     0.148    22.213 Allocation Failure Allocation Failure

C:\Users\as892333>jstat -gccapacity 30324
  NGCMN   NGCMX   NGC   S0C   S1C   EC   OGCMN   OGCMX   OGC   OC   MCMX   MC   CCSMN   CCSMX   CCSC   YGC   FGC   CGC   CGCT   GCT
  192.0  1107520.0  173376.0  17280.0  17280.0  138816.0  64.0  7209408.0  5808688.0  5808688.0  0.0  1056768.0  4864.0  0.0  1048576.0  512.0  612   0   0.269  35.689

C:\Users\as892333>jstat -gc 30324
  S0C   S1C   S0U   S1U   EC   EU   OC   OU   MC   MU   CCSC   CCSU   YGC   YGCT    FGC     FGCT    CGC     CGCT     GCT
  17280.0  17280.0  0.0  17279.1  138816.0  0.0  5808688.0  2404206.3  4864.0  3910.2  512.0  342.5  798    35.420   0.000   0.000   44     0.269  35.689

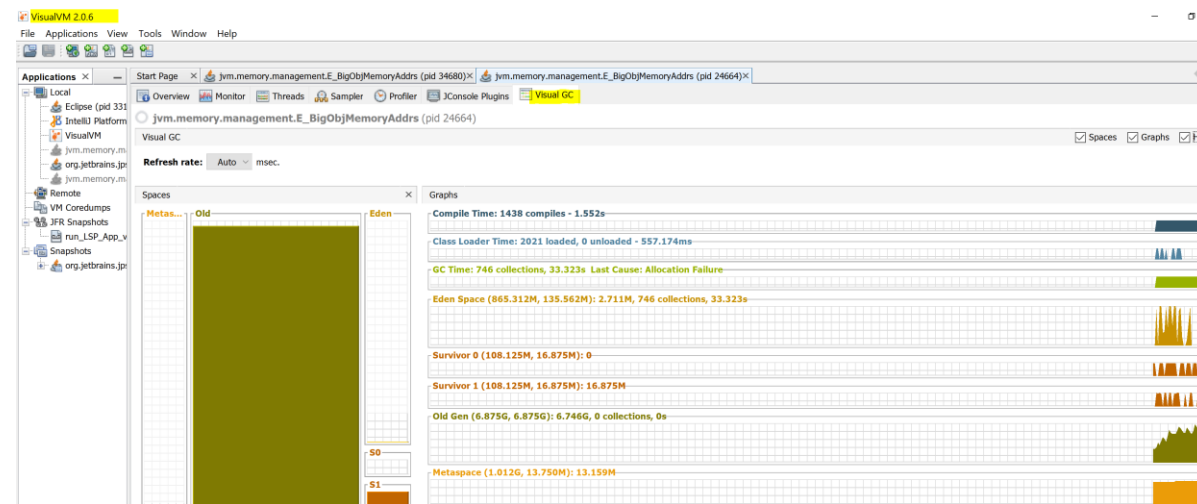
C:\Users\as892333>jstat -gc 30324 1000 10
  S0C   S1C   S0U   S1U   EC   EU   OC   OU   MC   MU   CCSC   CCSU   YGC   YGCT    FGC     FGCT    CGC     CGCT     GCT
  17280.0  17280.0  17280.0  0.0  138816.0  97152.2  5808688.0  2666411.2  4864.0  3910.2  512.0  342.5  873    39.133   0.000   0.000   48     0.289  39.422
  17280.0  17280.0  17280.0  0.0  138816.0  69394.5  5808688.0  3097911.0  4864.0  3910.2  512.0  342.5  887    39.715   0.000   0.000   50     0.305  40.020
  17280.0  17280.0  17280.0  0.0  138816.0  66618.7  5808688.0  2740607.6  4864.0  3910.2  512.0  342.5  901    40.334   0.000   0.000   50     0.305  40.639
  17280.0  17280.0  17280.0  0.0  138816.0  127685.7  5808688.0  2907835.0  4864.0  3910.2  512.0  342.5  915    40.976   0.000   0.000   51     0.305  41.281
  17280.0  17280.0  17280.0  0.0  138816.0  0.0  5808688.0  2706199.8  4864.0  3910.2  512.0  342.5  928    41.619   0.000   0.000   52     0.318  41.937
  17280.0  17280.0  0.0  17279.1  138816.0  0.0  5808688.0  2474330.0  4864.0  3910.2  512.0  342.5  939    42.244   0.000   0.000   52     0.318  42.562
  17280.0  17280.0  0.0  17279.1  138816.0  33309.4  5808688.0  3128731.5  4864.0  3910.2  512.0  342.5  952    42.831   0.000   0.000   53     0.319  43.150
  17280.0  17280.0  0.0  17279.1  138816.0  38861.0  5808688.0  3108791.6  4864.0  3910.2  512.0  342.5  965    43.461   0.000   0.000   54     0.335  43.796
  17280.0  17280.0  17280.0  0.0  138816.0  63842.8  5808688.0  2837968.4  4864.0  3910.2  512.0  342.5  979    44.070   0.000   0.000   54     0.335  44.404
  17280.0  17280.0  0.0  17279.1  138816.0  0.0  5808688.0  3146669.0  4864.0  3910.2  512.0  342.5  991    44.702   0.000   0.000   56     0.345  45.047
```

CODE: E_BigObjMemoryAdrrs

- **Visualvm** - VisualVM monitors and troubleshoots applications running on Java

Install Pluin for VisualGC

- **visualgc** – get more info in VM what is happening on GC





THANK YOU

References

https://docs.oracle.com/javase/8/docs/technotes/guides/vm/gctuning/g1_gc_tuning.html

<https://www.baeldung.com/jvm-experimental-garbage-collectors>

<https://www.journaldev.com/2856/java-jvm-memory-model-memory-management-in-java>

<https://docs.oracle.com/javase/7/docs/technotes/guides/vm/G1.html>

<http://tutorials.jenkov.com/java-concurrency/java-memory-model.html>

https://www.youtube.com/watch?v=WU_mqNBEacw

<https://shipilev.net/>

<https://www.youtube.com/watch?v=Gee7QfoY8ys> https://www.youtube.com/watch?v=WU_mqNBEacw